

Arithmetic Pipeline

32

Basic Concept

- Various arithmetic operations require a set of simpler computations to be carried out in sequence.
 - Can be split into stages with buffers in between, and run in a pipeline.
- Useful when several similar calculations are required to be carried out in sequence.
 - **Example:** Vector operations

```
for (i=0;i<64;i++)
    A[i] = B[i] * C[i];
```

33

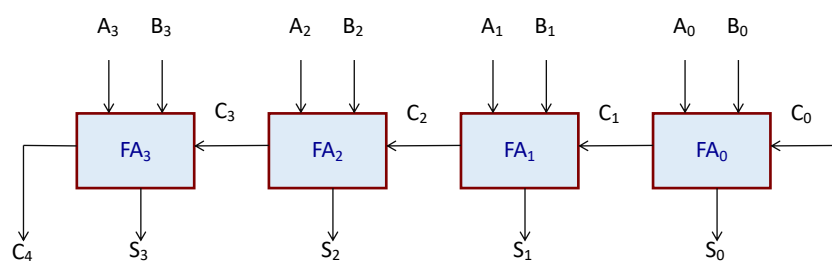
Fixed Point Addition Pipeline

- We have seen how a ripple-carry adder works.
 - Rippling of the carries gives it a bad worst-case performance.
- We explore whether pipelining can improve the performance.
- *Assumption*: delay of a latch is comparable to the delay of a full adder.

34

34

A 4-bit Ripple Carry Adder

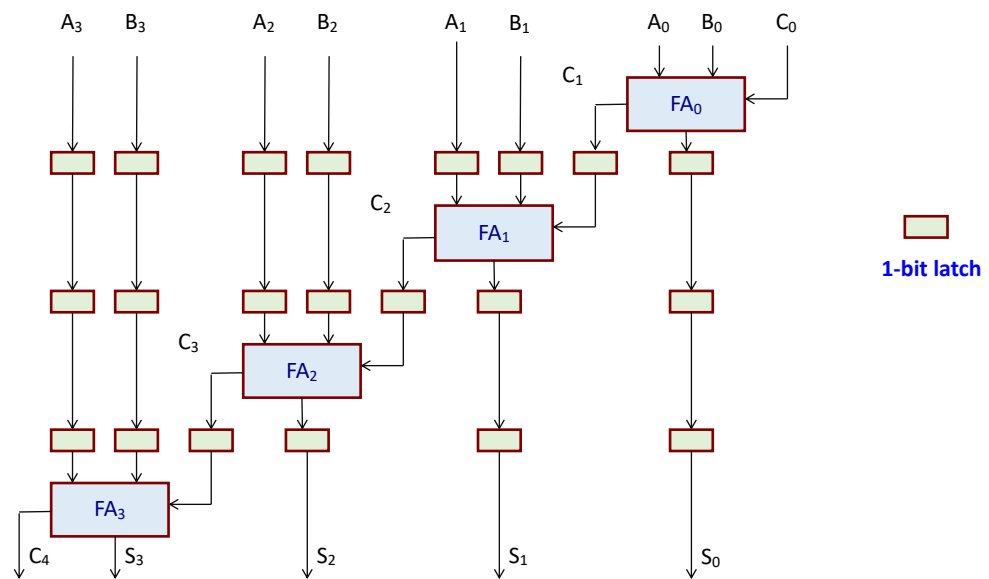


Worst-case delay $\approx 4 \times$ (carry generation time in FA)

35

35

4-bit pipelined ripple-carry adder



36

36

- Delay of a full adder = t_{FA}
- Delay of a 1-bit latch = t_L
- Clock period $T \geq (t_{FA} + t_L)$
- After the pipeline is full, one result (sum) is generated every time T .
 - Convenient for vector addition kind of applications.


```
for (i=0; i<10000; i++)
    a[i] = b[i] + c[i];
```

37

37

Fixed Point Multiplication Pipeline

- Combinational array multiplier using carry save adders can be directly implemented as a pipeline.
- The multiplier recoding logic will contain $8 \times 8 = 64$ AND gates to generate the partial product terms.



38

Floating-Point Addition

- Floating-point addition requires the following steps:
 - a) Compare exponents and align mantissas.
 - b) Add mantissas.
 - c) Normalize result.
 - d) Adjust exponent.
- Subtraction is similar.

Example:

$$A = 0.9504 \times 10^3$$

$$B = 0.8200 \times 10^2$$

Align mantissa: 0.0820

Add mantissa: $0.9504 + 0.0820 = 1.0324$

Normalize: 0.10324

Adjust exponent: $3 + 1 = 4$

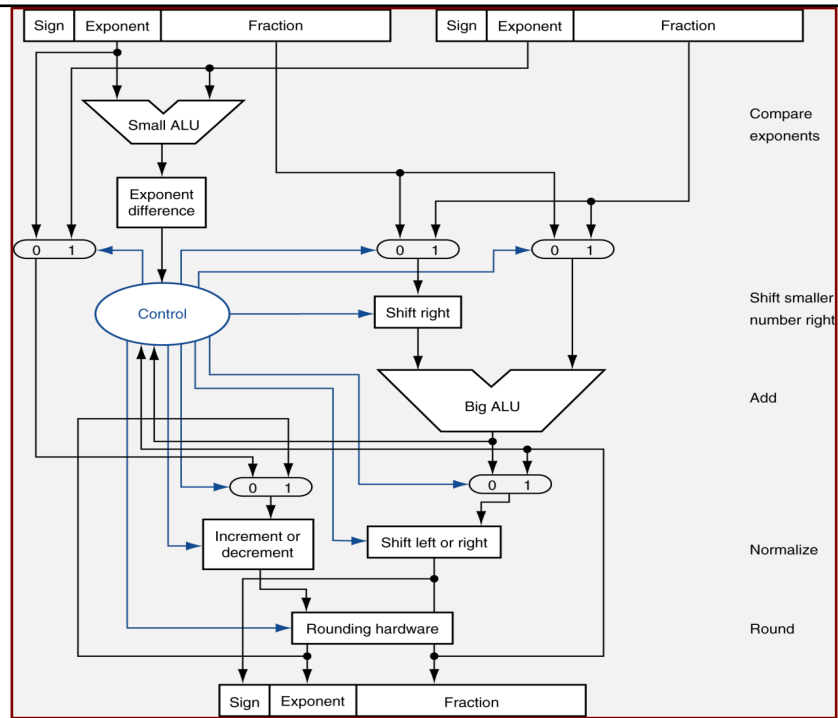
Sum = 0.10324×10^4

39

39

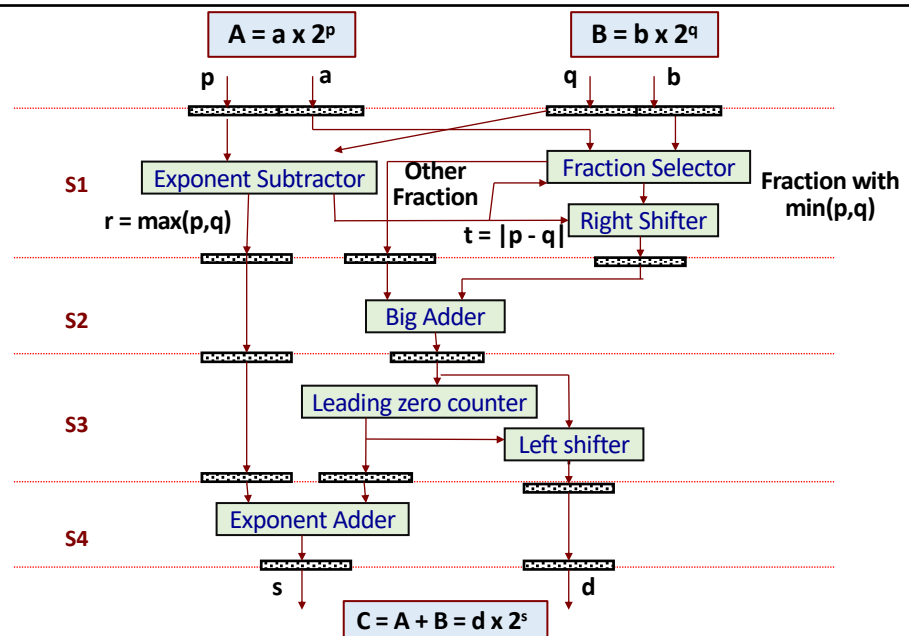
Floating-Point Addition Hardware

- The last step of rounding is required in IEEE-754 format.



40

4-Stage Floating Point Adder



41

Floating-Point Multiplication

- Floating-point multiplication requires the following steps:
 - Add exponents.
 - Multiply mantissas.
 - Normalize result.
- Division is similar.

A last step of rounding is required in IEEE-754 format.

Example:

$$A = 0.9504 \times 10^3$$

$$B = 0.8200 \times 10^2$$

Add exponents: $3 + 2 = 5$

Multiply mantissa: $0.9504 \times 0.8200 = 0.7793$

Normalize: 0.7793 (no change)

Product = 0.7793×10^5

42

42

Pipelining the MIPS32 Data Path

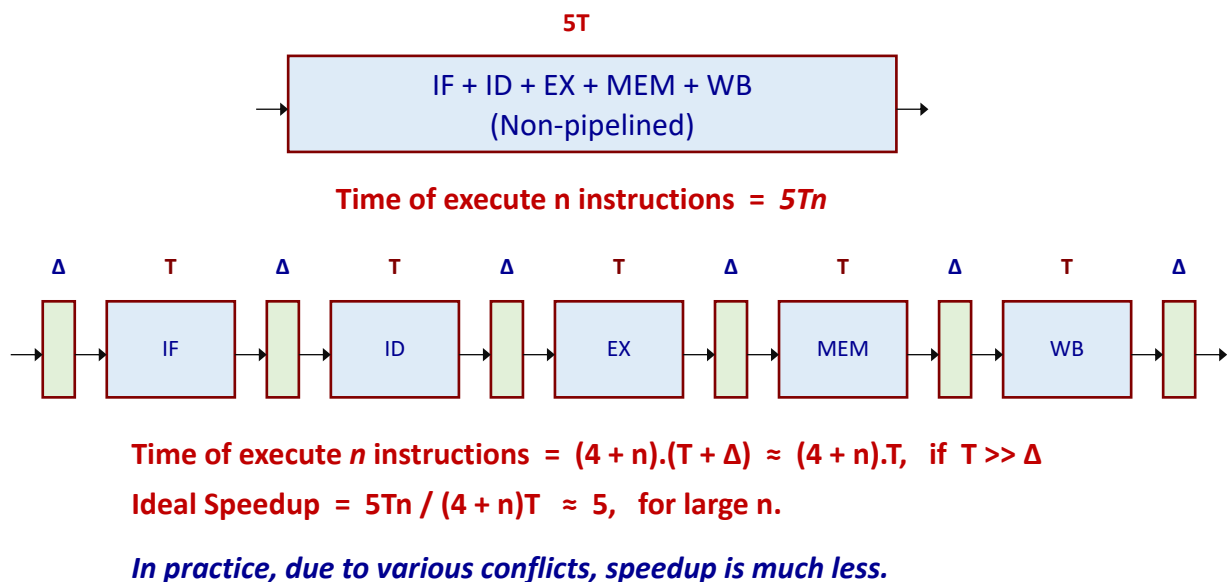
43

Introduction

- Basic requirements for pipelining the MIPS32 data path:
 - We should be able to start a new instruction every clock cycle.
 - Each of the five steps mentioned before (IF, ID, EX, MEM and WB) becomes a pipeline stage.
 - Each stage must finish its execution within one clock cycle.
- Since execution of several instructions are overlapped, we must ensure that there is no conflict during the execution.
 - Simplicity of the MIPS32 instruction set makes this evaluation quite easy.
 - We shall discuss these issues in some detail.

44

44



45

45

Clock Cycles								
Instruction	1	2	3	4	5	6	7	8
i	IF	ID	EX	MEM	WB			
$i+1$		IF	ID	EX	MEM	WB		
$i+2$			IF	ID	EX	MEM	WB	
$i+3$				IF	ID	EX	MEM	WB

Instr- i finishes
Instr- $(i+1)$ finishes
Instr- $(i+2)$ finishes
Instr- $(i+3)$ finishes

46

Clock Cycles								
Instruction	1	2	3	4	5	6	7	8
i	IF	ID	EX	MEM	WB			
$i+1$		IF	ID	EX	MEM	WB		
$i+2$			IF	ID	EX	MEM	WB	
$i+3$				IF	ID	EX	MEM	WB

Some examples of conflict:

- IF & MEM: In clock cycle 4, both instructions i and $i+3$ access memory.
 - Solution: use separate instructions and data cache.*
- ID & WB: In clock cycle 5, both instructions i and $i+3$ access register bank.
 - Solution: allow both read and write access to registers in the same clock cycle.*

47

Advantages of Pipelining


- In the non-pipelined version, the execution time of an instruction is equal to the combined delay of the five stages (say, $5T$).
- In the pipelined version, once the pipeline is full, one instruction gets executed after every T time.
 - Assuming all state delays are equal (equal to T), and neglecting latch delay.
- However, due to various conflicts between instructions (called *hazards*), we cannot achieve the ideal performance.
 - Several techniques have been proposed to improve the performance.
 - To be discussed.

48

48

Some Observations

- a) To support overlapped execution, peak memory bandwidth must be increased *5 times* over that required for the non-pipelined version.
- An instruction fetch occurs every clock cycle.
 - Also there can be two memory accesses per clock cycle (one for instruction and one for data).
 - Separate instruction and data caches are typically used to support this.




49

49

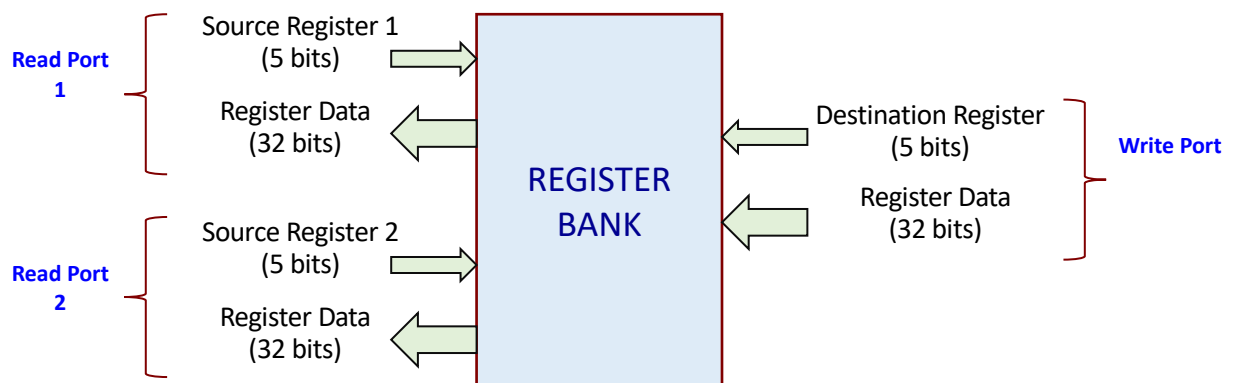
b) The register bank is accessed both in the stages ID and WB.

- ID requires 2 register reads, and WB requires 1 register write.
- We thus have the requirement of *2 reads and 1 write* in every clock cycle.
- Two register reads can be supported by having two register read ports.
- Simultaneous reads and write may result in clashes (e.g., same register used).
 - Solution adopted in MIPS32 pipeline is to perform the write during the *first half* of the clock cycle, and the reads during the *second half* of the clock cycle.



50

50



51

51

- c) Since a new instruction is fetched every clock cycle, it is required to increment the **PC** on each clock.
- PC updating has to be done *during IF stage itself*, as otherwise the next instruction cannot be fetched.
 - In the non-pipelined version discussed earlier, this was done during the MEM stage.

52

52

Basic Performance Issues in a Pipeline



- Register stages are inserted between pipeline stages, which increases the execution time of an individual instruction.
 - Because of overlapped execution of instructions, throughput increases.
- The clock period ***T*** has to be chosen suitably based on:
 - a) Slowest stage in the pipeline.
 - b) Clock skew and jitter.
 - c) Register setup time: minimum time the register input must be held stable before the active clock edge arrives.

53

53

Example 1

- Consider the 5-stage MIPS32 pipeline, with the following features:
 - Pipeline clock rate of 1GHz (i.e. 1 ns clock cycle time).
 - For a non-pipelined implementation, ALU operations and branches take 4 cycles, while memory operations take 5 cycles.
 - Relative frequencies of ALU operations, branches and memory operations are 50%, 15%, and 35% respectively.
 - In the pipelined implementation, due to clock skew and setup time, the clock cycle time increases by 0.25 ns.
 - Calculate the estimated speedup of the pipelined implementation.

54

54

• Solution:

- For non-pipelined processor:
 - Average instruction execution time = Clock cycle time x Average CPI

$$= 1 \text{ ns} \times (0.50 \times 4 + 0.15 \times 4 + 0.35 \times 5) = 4.35 \text{ ns}$$
- For pipelined processor:
 - Clock cycle time = $1 + 0.25 = 1.25 \text{ ns}$
 - In the steady state, one instruction will get executed every clock cycle.
 - Speedup = $4.35 / 1.25 = 3.48$

55

55

Micro-operations for Non-pipelined MIPS32

IF

```
IR ← Mem [PC];
NPC ← PC + 4;
```

ID

```
A ← Reg [rs];
B ← Reg [rt];
Imm ← (IR15)16 ## IR15..0
Imm1 ← IR25..0 ## 00
```

EX

```
ALUOut ← A + Imm;
(a) Memory
```

```
ALUOut ← A func B;
(b) R-R ALU
```

```
ALUOut ← A func Imm;
(c) R-IMM ALU
```

```
ALUOut ← NPC + (Imm << 2);
cond ← (A op 0);
(d) Branch
```

MEM

```
PC ← NPC;
LMD ← Mem [ALUOut];
(a) Load
```

```
PC ← NPC;
Mem [ALUOut] ← B;
(b) Store
```

```
if (cond) PC ← ALUOut;
else PC ← NPC;
(c) Branch
```

```
PC ← NPC;
(d) Others
```

56

56

WB

```
Reg [rd] ← ALUOut;
```

(a) R-R ALU

```
Reg [rt] ← ALUOut;
```

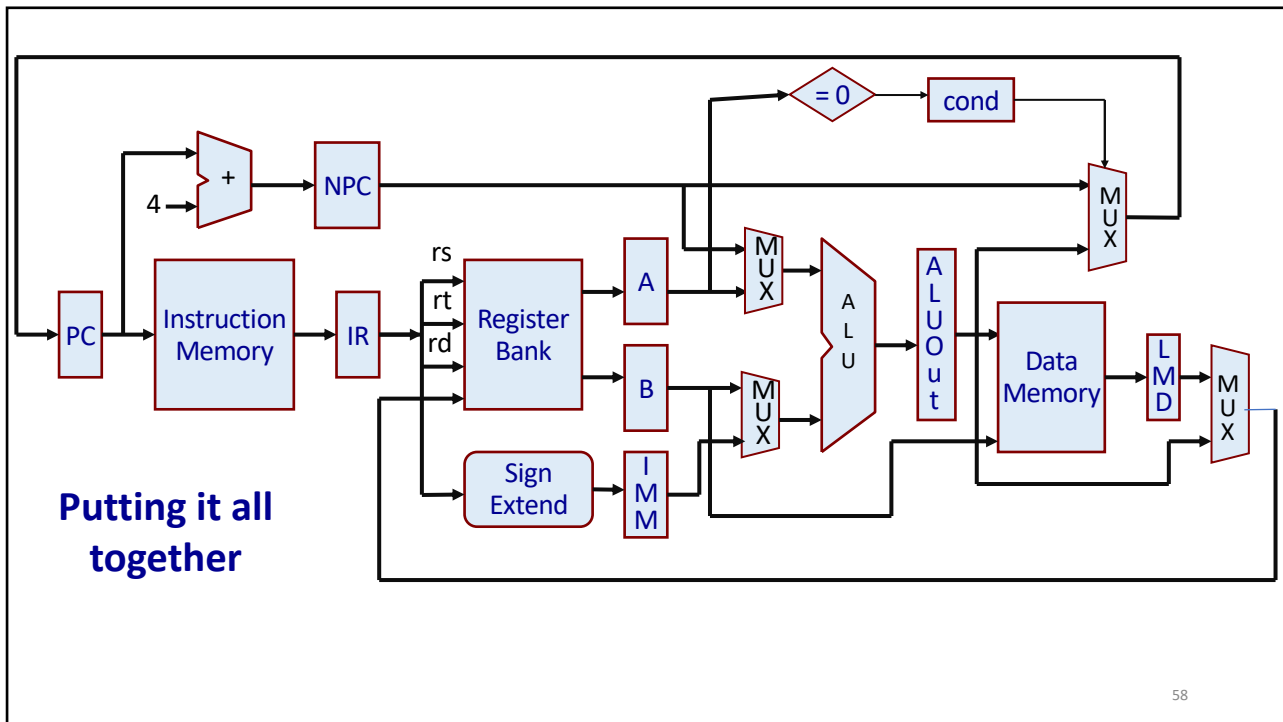
(b) R-IMM ALU

```
Reg [rt] ← LMD;
```

(c) Load

57

57



58

Micro-operations for Pipelined MIPS32

- Convention used:

- Many of the temporary registers required in the data path are included as part of the inter-stage latches.
- IF/ID: denotes the latch stage between the IF and ID stages.
- ID/EX: denotes the latch stage between the ID and EX stages.
- EX/MEM: denotes the latch stage between the EX and MEM stages.
- MEM/WB: denotes the latch stage between the MEM and WB stages.

- **Example:**

- ID/EX.A refers to a register **A** that is implemented as part of the **ID/EX** latch stage.

59

59

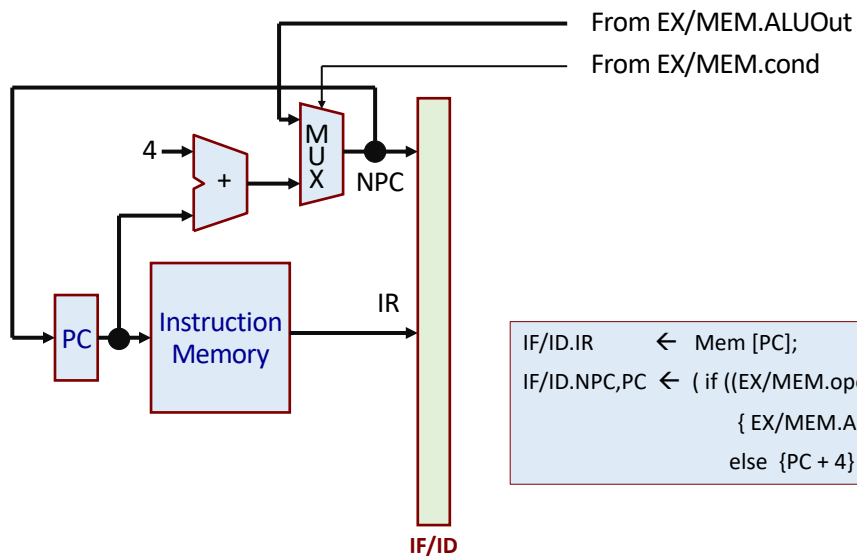
(a) Micro-operations for Pipeline Stage IF

```

IF/ID.IR    ← Mem [PC];
IF/ID.NPC,PC ← ( if ((EX/MEM.opcode == branch) & EX/MEM.cond)
                { EX/MEM.ALUOut}
                else {PC + 4} );
  
```

60

60



```

IF/ID.IR    ← Mem [PC];
IF/ID.NPC,PC ← ( if ((EX/MEM.opcode == branch) & EX/MEM.cond)
                { EX/MEM.ALUOut}
                else {PC + 4} );
  
```

61

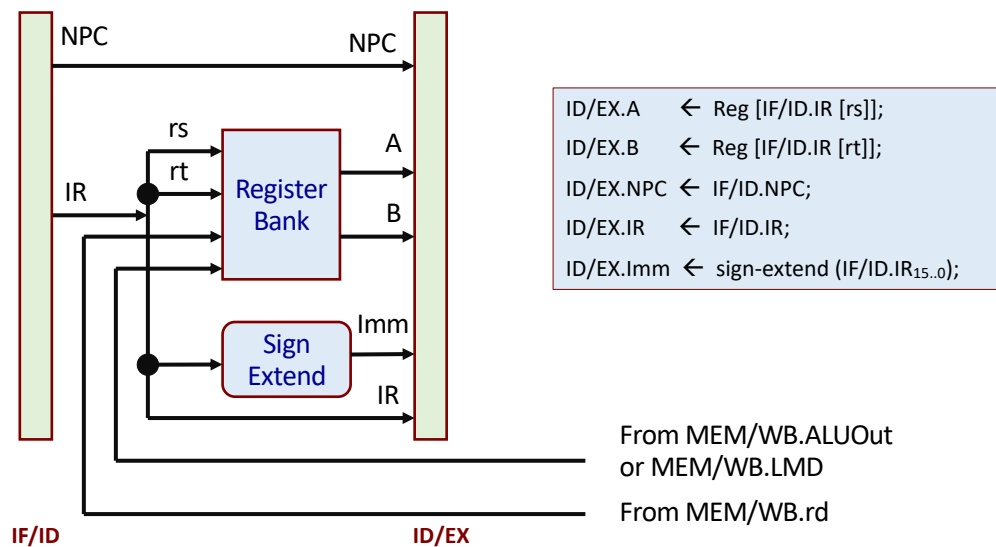
61

(b) Micro-operations for Pipeline Stage ID

$ID/EX.A \leftarrow \text{Reg} [IF/ID.IR [rs]];$
 $ID/EX.B \leftarrow \text{Reg} [IF/ID.IR [rt]];$
 $ID/EX.NPC \leftarrow IF/ID.NPC;$
 $ID/EX.IR \leftarrow IF/ID.IR;$
 $ID/EX.Imm \leftarrow \text{sign-extend} (IF/ID.IR_{15..0});$

62

62



63

63

(c) Micro-operations for Pipeline Stage EX

$$\text{EX/MEM.IR} \leftarrow \text{ID/EX.IR};$$

$$\text{EX/MEM.ALUOut} \leftarrow \text{ID/EX.A func ID/EX.B};$$

R-R ALU

$$\text{EX/MEM.ALUOut} \leftarrow \text{ID/EX.NPC} + (\text{ID/EX.Imm} \ll 2);$$

$$\text{EX/MEM.cond} \leftarrow (\text{ID/EX.A} == 0);$$

BRANCH

$$\text{EX/MEM.IR} \leftarrow \text{ID/EX.IR};$$

$$\text{EX/MEM.ALUOut} \leftarrow \text{ID/EX.A func ID/EX.Imm};$$

R-M ALU

$$\text{EX/MEM.IR} \leftarrow \text{ID/EX.IR};$$


$$\text{EX/MEM.ALUOut} \leftarrow \text{ID/EX.A} + \text{ID/EX.B};$$

$$\text{EX/MEM.B} \leftarrow \text{ID/EX.B};$$

LOAD / STORE

64

64



$$\text{EX/MEM.IR} \leftarrow \text{ID/EX.IR};$$

$$\text{EX/MEM.ALUOut} \leftarrow \text{ID/EX.A func ID/EX.B};$$

$$\text{EX/MEM.IR} \leftarrow \text{ID/EX.IR};$$

$$\text{EX/MEM.ALUOut} \leftarrow \text{ID/EX.A func ID/EX.Imm};$$

$$\text{EX/MEM.IR} \leftarrow \text{ID/EX.IR};$$

$$\text{EX/MEM.ALUOut} \leftarrow \text{ID/EX.A} + \text{ID/EX.B};$$

$$\text{EX/MEM.B} \leftarrow \text{ID/EX.B};$$

$$\text{EX/MEM.ALUOut} \leftarrow \text{ID/EX.NPC} + (\text{ID/EX.Imm} \ll 2);$$

$$\text{EX/MEM.cond} \leftarrow (\text{ID/EX.A} == 0);$$

65

65

(d) Micro-operations for Pipeline Stage MEM

MEM/WB.IR \leftarrow EX/MEM.IR;
MEM/WB.ALUOut \leftarrow EX/MEM.ALUOut;

ALU

MEM/WB.IR \leftarrow EX/MEM.IR;
MEM/WB.LMD \leftarrow Mem [EX/MEM.ALUOut];

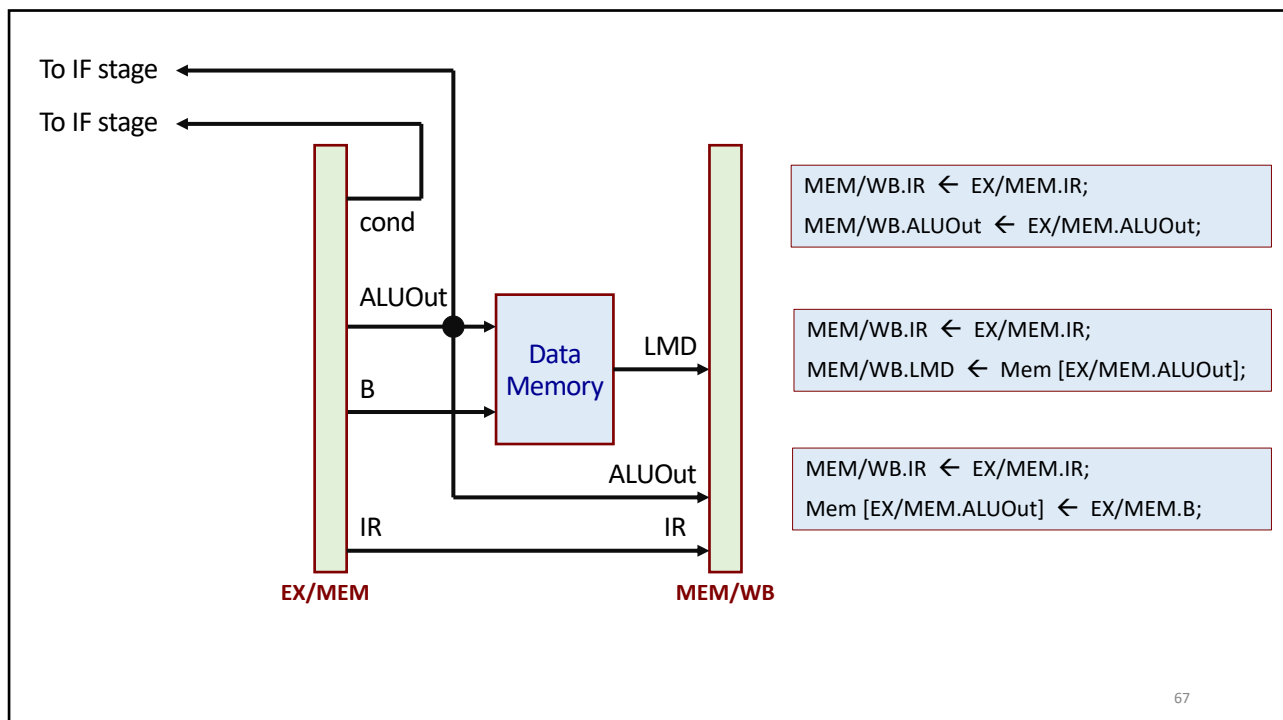
LOAD

MEM/WB.IR \leftarrow EX/MEM.IR;
Mem [EX/MEM.ALUOut] \leftarrow EX/MEM.B;

STORE

66

66



67

67

(e) Micro-operations for Pipeline Stage WB

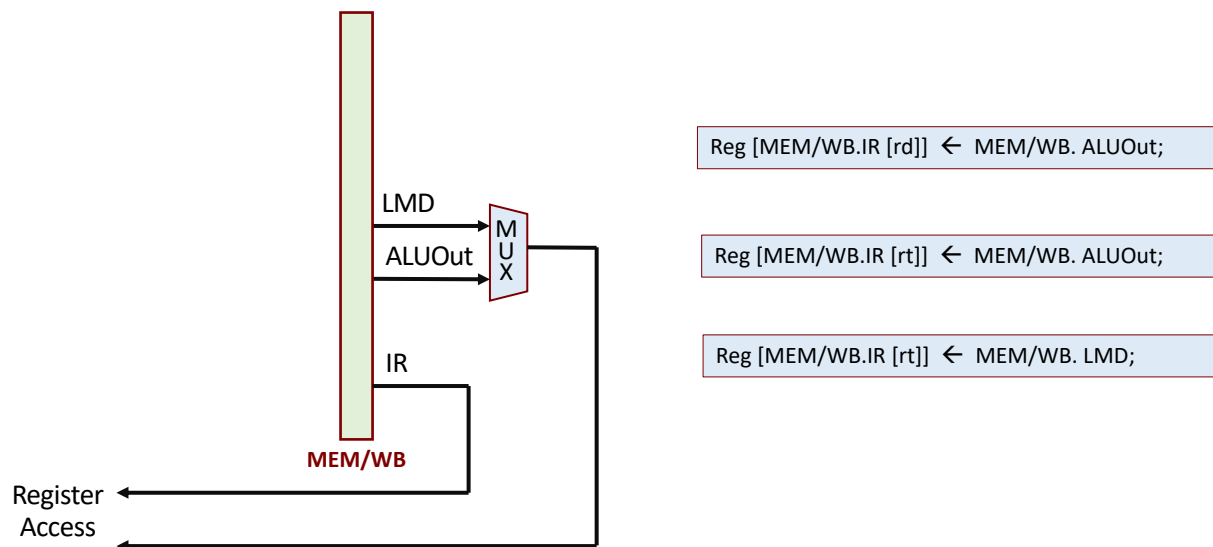
Reg [MEM/WB.IR [rd]] \leftarrow MEM/WB. ALUOut; **R-R ALU**

Reg [MEM/WB.IR [rt]] \leftarrow MEM/WB. ALUOut; **R-M ALU**

Reg [MEM/WB.IR [rt]] \leftarrow MEM/WB. LMD; **LOAD**

68

68



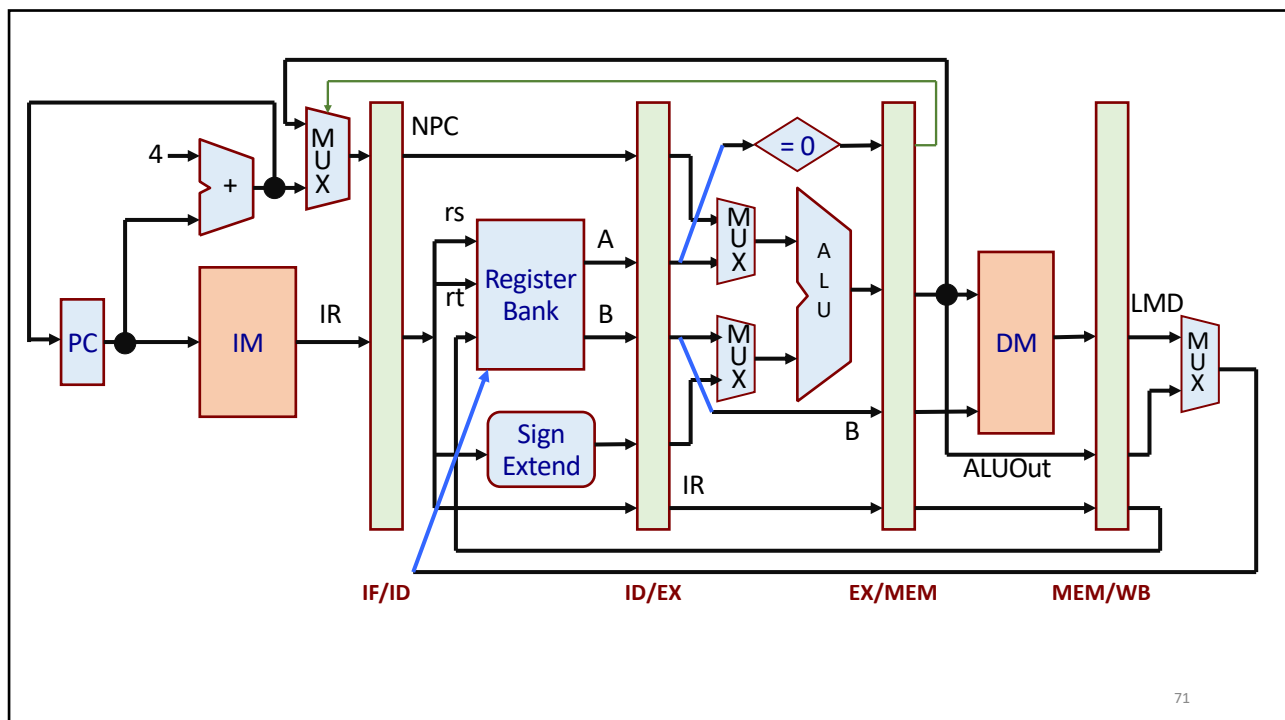
69

69

PUTTING IT ALL TOGETHER :: MIPS32 PIPELINE

70

70



71

71